

HW 3 - Report

Table of Contents

1.	Design Discussion.....	2
1.1	Pre-Processing Job	2
1.2	Page Rank Job	4
1.3	Top-K Job.....	7
1.4	Data Transferred	8
2.	Performance Comparison	10
2.1	Running Time.....	10
2.2	Running Time Comparison	11
2.3	Top-100 Wikipedia Pages.....	12
3.	Convergence Estimation	17

1. Design Discussion

1.1 Pre-Processing Job

I have followed the same per-processing approach as described in the Assignment. I have used given SAX XML Parser. I have modified three things in the given parser:

1. Do not include self links
2. Removed the duplicate nodes from the adjacency list
3. Fixed the “ & ” issue, some pages was not parsed due to this issue.

I have created my own input file which covers all cases required in assignment and calculated manual page ranks for this input for 10 iterations. And then compared each output of my all programs to the manual calculation and they are the same.

PreProcessingJob emits the nodeId and its adjacentList for all successfully parsed documents. And the job returns the total number of pages from global counter.

Pseudo-Code

```
class Mapper
{
    map(Line line)
    {
        (nodeId, adjacencyList) = parse(line);

        // Emit nodeId with its adjacencyList
        emit(nodeId, adjacencyList);

        // Emit each adjacentNode with empty adjacencyList
        // If some nodes are in adjacencyList but do not have document for them
        // then this dangling nodes need to be considered
        for(String adjacentNodeId : adjacencyList)
        {
            emit(adjacentNodeId, EMPTY_LIST);
        }
    }
}
```

```
class Reducer
{
    Long localNumberOfPages;

    setup()
    {
        localNumberOfPages = 0;
    }

    reduce(nodeId, [adjacencyList1, adjacencyList2,...])
    {
        adjacencyList = EMPTY_LIST;

        // get the first non empty adjacencyList if present
        // There can be multiple documents with same nodeId in input
        // so consider the first one only
        for each val in input list do
        {
            if(val is not EMPTY_LIST)
            {
                adjacencyList = val;
                break;
            }
        }

        // increment the number of pages
        localNumberOfPages++;

        // emit nodeId, its adjacencyList and its default page rank
        emit((nodeId, adjacencyList), NULL);
    }

    // cleanup increments the global counter by localNumberOfPages count.
    cleanup()
    {
        increment global counter with localNumberOfPages;
    }
}
```

On successful completion of above job we will return the total number of pages from global counters. Total number of pages should be same as number of reduce outputs.

1.2 Page Rank Job

For calculating page rank (computing dangling factor delta), I have used Solution – 2 (Merge computation of delta in previous reduce phase) approach as per Learning module 6 - Graph Algorithms. Below are the all three approaches for calculating delta and the reasons why I have choose solution 2.

Solution-1 – Add a separate phase to each iteration to compute delta

During an iteration, first execute a Map Reduce program that computes delta. This is a simple global aggregation job, summing up Page Rank values for all dangling nodes. Then pass the newly computed delta as a parameter to modified Map Reduce program that updates all Page Ranks using new formula with delta. The drawback of this approach is that, in each iteration we are reading all the records in map to calculate the delta in global aggregation job. So here we have an overhead of reading all the inputs for each iteration.

Solution-2: Merge computation of delta in previous reduce phase. (Used Approach)

In this approach, instead of computing delta in a separate job in the beginning of iteration (i+1), it could already be computed at the end of iteration i in reduce call. In this approach we don't need the extra reading of all input records as we need in above approach. Only problem with this approach is, the page rank emitted by the reducer is not the correct one so in last iteration we need to add one more map only job to correct the page ranks calculated by the last iteration. In our program, we have a PageRankCorrectionJob which does the same. So here we can argue that we need one extra map read of all the inputs per entire job, but this can also be eliminated. In our program, We just need TopK page ranks and we can correct the output emitted by the last iteration in TopK job. But this will give the correct page rank of only top K pages and it is possible that we might need other records in future. So for the safer side, I have provided PageRankCorrectionJob to run before TopKJob.

Solution-3 – Order Inversion

The order inversion pattern can be applied to make sure each Reducer receives the old PageRank values of all dangling nodes. Each reducer can then compute delta right before executing any of the normal reduce calls. This approach has the drawback same as what we have seen in other Order Inversion algorithms. We need to have one reduce task only so we cannot have reduce task granularity. We can overcome this issue by emitting multiple copies of same dangling node's page rank such that all reducers receive exact one copy of it. But this creates map output duplication issue. We can overcome this output duplication issue in some extent by in-mapper combining. So this approach with in-mapper combining to dangling factor output in each map task and emitting that combined delta to all reducer (multiple reduce tasks) can be a good solution. But as per our problem needs we required only top k page ranks and in our solution 2 approach we can do that without using PageRankCorrectionJob which makes Solution 2 better than other two approaches.

Pseudo Code:

```
// Map processes the node with id n.
// N stores node n's current PageRank and its adjacency list
map(nid n, node N)
{
    if( isfirstIteration )
    {
        N.pageRank =  $1/|V|$ 
    }
    else
    {
        delta = getDeltaFromConfiguration()
        N.pageRank =  $\alpha/|V| + (1-\alpha)(\text{delta}/|V| + \text{N.pageRank})$ 
    }

    // Pass along the graph structure
    emit(nid n, N)

    if(|N.adjacencyList| != 0)
    {
        //compute contributions to send along outgoing links
        p = N.pageRank / |N.adjacencyList|
        for all nid m in N.adjacencyList do
            emit(nid m, p)
    }
    else
    {
        emit("Delta", N.pageRank)
    }
}

// In code, I have used In-mapper combining to combine delta and also used Combiner to
// combine the map outputs

// Reduce receives the node object from node m and
// the PageRank contributions for all m's inlinks
reduce(nid m, [p1, p2,...])
{
    s=0
    M=NULL

    for all p in [p1,p2,...] do
        if isNode(p) then
```

```
        // The node object was found : recover graph structure
        M = p
    else
        // A PageRank contribution from an inlink was
        // found : add it to the running sum
        s += p1

    if (m == "Delta")
        incrementGlobalCounter("Delta", s)
    else
        M.pageRank = s
        emit(nid m, node M)
}
// After each iteration we will pass computed delta from global counters to next job's
configuration.
```

1.3 Top-K Job

There are two approaches to find the top-K elements. First is to sort the input data and then the K largest records can easily be selected from the sorted file. This is an efficient method for very large K. Second approach is to avoid sorting if we have smaller value of K. The idea is to scan the input only once and use in-mapper combining to keep track of the top-K records in each map task. A single reduce call receives these local top-k lists and merges them into final result.

In our assignment K=100 which is a very small number, So I have used the second approach and not sorting as we do not need to transfer all input records from map to reduce, read again in reduce and write from reduce to file.

Pseudo-Code:

Pseudo code is same as the pseudo code given in Module – 5 : Basic Algorithms.

[2.13: Top-K records – Second Approach – find the local TopK in each Map task and merge them into global TopK in one reduce task.]

Addition to the given code, I have implemented MyTreeMap to handle the different pages with same page rank issue.

1.4 Data Transferred

Below table shows the amount of data transferred from Mappers to Reducers, and from Reducers to S3 (on AWS, job has configured to write on S3 instead of HDFS) in each phase of the entire program execution and in each iteration of page rank job. Data from syslog file of AWS EMR run of Wikipedia-full-html input execution on 11 machines (1 master and 10 workers).

Phase - Iteration	S3: Number of bytes read (S3 to Mappers)	Map output bytes	Reduce shuffle bytes (Combiner/Mapper to Reducer)	S3: Number of bytes written (Reducer to S3)
Pre-Pro.	7091101822	2088430787	1008848757	1057795115
1	1057795115	3210315684	1217336429	1180704724
2	1180704724	3214348041	1317781918	1182955026
3	1182955026	3212780522	1317996064	1181970571
4	1181970571	3213168630	1318206121	1183014094
5	1183014094	3213382361	1318341266	1183023055
6	1183023055	3213765772	1318386360	1183036262
7	1183036262	3212932819	1318206906	1181992025
8	1181992025	3212833986	1318236724	1181995559
9	1181995559	3214408272	1318455864	1183051088
10	1183051088	3213143682	1318350598	1181993998
Correction	1181993998	Map Only Job, Mapper to S3 : 142752032 bytes		
TopK	142752032	65692	59740	3211

Below table shows the number of records transferred between mapper, combiner and reducer. Data from syslog file of AWS EMR run of Wikipedia-full-html input execution on 11 machines (1 master and 10 workers).

Iteration	Map input records	Map output records	Combine input records	Combine output records	Reduce input records	Reduce output records
Pre-Pro.	7012253	54626996	54626996	19434887	19434887	3178227
1	3178227	54718394	54718394	19556653	19556653	3178227
2	3178227	54718394	54718394	19562914	19562914	3178227
3	3178227	54718394	54718394	19559839	19559839	3178227
4	3178227	54718394	54718394	19560635	19560635	3178227
5	3178227	54718394	54718394	19561348	19561348	3178227
6	3178227	54718394	54718394	19561719	19561719	3178227
7	3178227	54718394	54718394	19560157	19560157	3178227
8	3178227	54718394	54718394	19560043	19560043	3178227
9	3178227	54718394	54718394	19562772	19562772	3178227
10	3178227	54718394	54718394	19560535	19560535	3178227
Correction	3178227	3178227	Map Only Job			
TopK	3178227	1900	0	0	1900	100

In pre-processing, the input document size is very high. But in pre-processing, we discard most of the data from document and end up to only graph representation (nodeId : Adjacency List). So output of the reducer to S3 is comparatively smaller.

In first iteration, in input file (S3 to mapper), we have only nodeId and adjacencyList, but after processing the first iteration of page rank, reducers emit the page rank of each page along with its nodeId and adjacencyList. This increases the number of bytes written to S3 from reducers for first iteration.

For all other iteration, there is no bigger difference in amount of data transferred in each step, as we are processing & passing the kind of same data and just modifying the page rank values. In each iteration, amount of data transferred from Mapper to Reducer increases because we are passing the page rank contribution to each adjacent node along with node structure entries. Though Combiner helped a lot to reduce the reduce shuffle bytes from original Map output bytes. Reducer combines the map outputs and writes it to S3 so the number of bytes reduces again.

Last Correction job writes only NodeId and PageRank values amount of data transferred to S3 from reducer drastically reduced same for TopK(only K=100 records written) job as well.

2. Performance Comparison

2.1 Running Time

I have executed program on AWS EMR for both the inputs. I have used $\alpha = 0.15$ for my page rank calculations (Source: Wikipedia page of page rank algorithm). Though the value of α is not hardcoded and you can configure in Makefile.

Below is the screenshot of AWS EMR successful run:

	Name	ID	Status	Creation time (UTC-5)	Elapsed time
<input type="checkbox"/>	HW3 Cluster	j-2KC0ER06JRMSG	Terminated All steps completed	2017-02-25 22:39 (UTC-5)	34 minutes
<input type="checkbox"/>	HW3 Cluster	j-J2H4T3STQ6L	Terminated All steps completed	2017-02-25 20:56 (UTC-5)	51 minutes
<input type="checkbox"/>	HW3 Cluster	j-U2Y0CQ1HHKSE	Terminated All steps completed	2017-02-25 18:32 (UTC-5)	16 minutes

1st row: Wikipedia-full-html input run on 11 m4.large machines (1 master and 10 workers)

2nd row: Wikipedia-full-html input run on 6 m4.large machines (1 master and 5 workers)

3rd row: Wikipedia-simple-html input run on 6 m4.large machines (1 master and 5 workers)

Below is the running time of full Wikipedia input on both configurations:

Job	(Time in Seconds)	
	6 m4.large Machine	11 m4.large Machines
Pre-Processing Time	1021	631
Time to Run ten Iterations Of Page Rank	1532	947
Time to Find The Top -100 Pages	40	29

2.2 Running Time Comparison

As we see in running times, all the computation phases shown good speed up by increasing machines as expected. With higher number of machines, our work load is divided among them and due to parallel work overall time is reduced.

If we compare the speed up of all computational phases then first two phases (PreProcessing and PageRank) has shown higher speed up then the last phase (TopK). The running time decreased to 38% in first two phases while it reduced 27% in last phase. I have also expected the same as in first two phases our MapReduce framework utilizes all the machines whereas in TopK job we are creating only on reduce task so we cannot get advantage of increase in machines in reduce work of TopK job.

Number of reduce task depends on the available nodes to MapReduce framework. As per hadoop documentation *"The right number of reduces seems to be 0.95 or 1.75 multiplied by (<no. of nodes> * <no. of maximum containers per node>)." If we have higher number of machines then it will create more reduce task and we will get good parallelism which will reduce the time taken and increase the performance.*

In our case for first two phases, for 11 machines, total number of reduce task is 20 whereas for 6 machines they are reduced to 10. So in first two phases, higher machine gave better performance as the load is distributed among them nicely. For last phase, number of reduce tasks are independent of machines as we are setting it manually to 1. So here we are not getting advantage of increase in machines in reduce phase and it shows lesser speed up then above two phases. Adding to

2.3 Top-100 Wikipedia Pages

Below is the top-100 Wikipedia pages with highest page rank along with their page values sorted from highest to lowest for both the simple and full datasets. I find them reasonable as I guess there will be higher in links to the pages like United_States_09d4 from the other pages which result into higher page rank of this page. And we can see other countries as well in the list this means, the other wiki pages will have links in their web page to the related country page and which will boost the page rank of the country pages.

No	Wikipedia-full-html-output-11-machines		Wikipedia-simple-html-output-6-machines	
	Page Name	Page Rank	Page Name	Page Rank
1	United_States_09d4	0.00262288934616 4360	United_States_09d4	0.00518900900027 2900
2	2006	0.00122849783940 2600	Wikimedia_Commons_7b57	0.00480676647470 8800
3	United_Kingdom_5ad7	0.00120313676095 6090	Country	0.00394028468771 2630
4	Biography	0.00098207155818 9593	England	0.00275248143611 0630
5	2005	0.00091705986394 1628	Water	0.00268780962344 6530
6	England	0.00088020713412 7707	Animal	0.00255408756514 9160
7	Canada	0.00085590507934 7993	City	0.00251082408078 2450
8	Geo-graphic_coordinate_system	0.00077172622988 9266	United_Kingdom_5ad7	0.00235864709361 2220
9	France	0.00072502418722 5341	Germany	0.00235040169771 1510
10	2004	0.00071989680727 9462	Earth	0.00232473485995 4630
11	Australia	0.00068047669529 2216	France	0.00232360794714 2090
12	Germany	0.00065434528014 1820	Europe	0.00203809703716 7730
13	2003	0.00058738697782 8930	Wiktionary	0.00175388421427 6070
14	India	0.00058341977400 5353	English_language	0.00174967712175 4420
15	Japan	0.00058285604737 3971	Government	0.00173234465210 3290
16	Inter-net_Movie_Database_7ea7	0.00053350696600 3842	Computer	0.00171684048471 3360
17	Europe	0.00050926763789	India	0.00171317091838

		9670		4910
18	Record_label	0.00049145956758 4876	Money	0.00166738369802 2790
19	2001	0.00048701215830 3275	Japan	0.00155169056853 5440
20	2002	0.00048286324681 9372	Plant	0.00152355950935 9920
21	World_War_II_d045	0.00047805043674 7285	Italy	0.00150743309049 7990
22	Population_density	0.00047034299613 8859	Canada	0.00148140734345 2890
23	Music_genre	0.00046721175179 6516	Spain	0.00147112369222 3520
24	2000	0.00046466103240 1972	Food	0.00142468684896 7640
25	Italy	0.00044579274603 5966	Human	0.00141209700626 9630
26	Wiktionary	0.00043620978702 9988	China	0.00139671506127 2910
27	Wikimedia_Commons_7b57	0.00043529472397 9010	People	0.00138224852505 5760
28	London	0.00043480265906 0111	Australia	0.00132985424075 0500
29	English_language	0.00041850352295 4299	Asia	0.00128443617113 6100
30	1999	0.00040593698578 7724	Capital_(city)	0.00127426842125 1940
31	Spain	0.00036295379192 9990	Television	0.00126499722576 0360
32	1998	0.00035631063673 4167	Sun	0.00126021008117 8010
33	Russia	0.00034390662497 8832	Number	0.00124323622892 8820
34	1997	0.00033728202493 3117	State	0.00124037568145 4610
35	Television	0.00033629712380 0393	Sound	0.00123521166722 1920
36	New_York_City_1428	0.00033462877684 3839	Science	0.00123254317535 9430
37	Football_(soccer)	0.00032614642134 4763	Mathematics	0.00123105663929 5570
38	1996	0.00032362786609 1999	Metal	0.00119230462374 9440
39	Census	0.00032355337188 6843	Year	0.00117709258351 0610
40	Scotland	0.00032218915558 6006	2004	0.00117335731376 8480

41	1995	0.00031015464104 7029	Language	0.00115016588485 7740
42	China	0.00030864301289 6061	Russia	0.00114618177921 2580
43	Population	0.00030432048264 0489	Wikipedia	0.00112333028098 8190
44	Square_mile	0.00030405610072 0908	Religion	0.00109856669996 6040
45	Scientific_classification	0.00030401196802 8030	19th_century	0.00109653914178 0080
46	California	0.00030166740429 4186	Music	0.00108743132321 4420
47	1994	0.00029069115857 5429	Scotland	0.00105480073500 6320
48	Sweden	0.00028762080340 1212	20th_century	0.00105370498325 8870
49	Public_domain	0.00028741610678 1891	Greece	0.00104922273293 4630
50	Film	0.00028626913927 5013	Latin	0.00102986061318 7450
51	Record_producer	0.00028411019015 9314	London	0.00102735544285 1320
52	New_Zealand_2311	0.00028310205424 5546	Greek_language	0.00100435725665 0290
53	New_York_3da4	0.00027888263555 7677	Energy	0.00099901181037 9402
54	Netherlands	0.00027667367831 8118	World	0.00098635084799 7666
55	Marriage	0.00027581329226 8387	Centuries	0.00097590586513 6575
56	1993	0.00027482489872 0027	Culture	0.00094520396521 1297
57	United_States_Census_Bureau_2c85	0.00027466710694 8521	History	0.00093646960342 5431
58	1991	0.00027189525612 9761	Liquid	0.00091452309680 0025
59	1990	0.00026832611879 8734	Netherlands	0.00090572450764 8975
60	1992	0.00026636924637 7862	Planet	0.00090493226223 9007
61	Politician	0.00026489490763 6039	Light	0.00090167635268 6388
62	Album	0.00026056445391 8753	Society	0.00090149206214 5202
63	Latin	0.00026045636147 8826	Atom	0.00089002264065 2958
64	Actor	0.00025833956084	Wikime-	0.00088844007077

		6043	dia_Foundation_83d9	6097
65	Ireland	0.00025810638770 3505	Scientist	0.00088838361057 3494
66	Per_capita_income	0.00025564272454 9733	Image	0.00088768848602 1999
67	Studio_album	0.00025186026160 0276	Law	0.00088629080559 8419
68	Poverty_line	0.00025116500863 4912	Geography	0.00087884516145 4898
69	Km²	0.00024950659722 5139	List_of_decades	0.00087857429428 3701
70	1989	0.00024689373252 3362	Uni- form_Resource_Locator_1b 4e	0.00086188450636 3216
71	Norway	0.00024092192871 2262	Africa	0.00086056996715 2458
72	Website	0.00023901785842 5252	Turkey	0.00084488636788 9011
73	1980	0.00023532199065 3687	Inhabitant	0.00083047948823 2369
74	Animal	0.00022937863976 8423	Capital_city	0.00082304881404 3760
75	Area	0.00022920870296 2197	Plural	0.00082151559551 0230
76	1986	0.00022703304683 4427	Electricity	0.00081372300166 6497
77	Personal_name	0.00022623653898 7703	Poland	0.00079723790431 5345
78	Poland	0.00022611544290 5689	Building	0.00079712389257 2039
79	Brazil	0.00022568426663 7636	Car	0.00079465406062 3909
80	1985	0.00022402906904 2517	Sweden	0.00079171255623 4121
81	1987	0.00022330540325 8049	Book	0.00079148847053 1959
82	1983	0.00022175338663 6217	Biology	0.00078693289643 1402
83	1982	0.00022109653273 6380	War	0.00077081729454 8049
84	French_language	0.00021938452665 4706	Chemical_element	0.00076816079591 9691
85	1981	0.00021934801195 2457	God	0.00076093572189 1388
86	1979	0.00021932859347 5849	North_America_e7c4	0.00075628686441 6695
87	1984	0.00021879019421	September_7	0.00075477818126

		2315		4075
88	World_War_I_9429	0.00021869239369 5080	Website	0.00074629735006 0417
89	1988	0.00021857418021 8445	Nation	0.00074266715264 0605
90	Paris	0.00021801968011 5686	Politics	0.00073971037875 8902
91	1974	0.00021797486466 4915	2006	0.00073329001722 5940
92	Mexico	0.00021567359642 7163	Fish	0.00073223711129 0969
93	19th_century	0.00021185635773 6199	Species	0.00073087111762 9328
94	1970	0.00021132429567 9975	Mammal	0.00072167441359 4913
95	January_1	0.00021086426991 0462	Island	0.00071780902030 3589
96	USA_f75d	0.00021070868243 2733	Portugal	0.00071710705966 0586
97	1975	0.00020860183512 0045	Gas	0.00071555153665 3913
98	1976	0.00020846744997 3305	River	0.00071157775130 0913
99	Africa	0.00020779805256 2159	Switzerland	0.00070610750743 8502
10 0	South_Africa_1287	0.00020736244197 3637	World_War_II_d045	0.00070203049315 8166

3. Convergence Estimation

I have coded to measure the convergence achieved after each iteration of page rank calculation. In each iteration, program calculates the absolute difference in old and new page rank of each node and sums it up. And at the end of iteration, the sum is divided by the number of total pages to get the average difference per node and logs it to log file.

Below is the convergence output from logs of AWS EMR execution of Wikipedia-full-html input with both configurations:

Iteration	6 m4.large Machine	11 m4.large Machines
1	0.000000184996236977	0.000000184996609594
2	0.000000076656124798	0.000000076656300012
3	0.000000030632383622	0.000000030632328128
4	0.000000013365304956	0.000000013365282971
5	0.000000006365162230	0.000000006365152615
6	0.000000003233759976	0.000000003233754581
7	0.000000001721090299	0.000000001721086772
8	0.000000000951050433	0.000000000951048336
9	0.000000000543148803	0.000000000543147593

(Note: It doesn't calculate for last iteration, as last page rank is corrected in PageRankCorrectionJob)

As we can see in above diagram, average difference in page rank in each iteration reduces which means page ranks are converging.

If we need the convergence up to specific decimal point then we can count the number of nodes whose page rank change is higher than specified margin in each iteration. This count will decrease in each iteration and when it reaches to zero, we will stop calculating the page rank.